

Slide 10: Data Engineers - AI-Ready Data Infrastructure

Objective: Define data engineering skills needed to build and maintain data infrastructure that supports AI/ML workloads.

Layout Composition:

The slide is a flex-col with a header sized to its content, a main body that grows, and a footer sized to its content. The main body is a grid with two columns (3fr 2fr). The left column is a flex-col with two stacked skill cards. The right column is a flex-col with a data pipeline visualization component at top and a tools card below.

Content Breakdown:

- Block 1:
 - Block Type: Text
 - Placement: Header section
 - Component Schema: “Role Header”
 - Content:
 - Role_Title: “Data Engineers”
 - Role_Icon: “Database”
 - Subtitle: “AI-Ready Data Infrastructure & Pipeline Engineering”
 - Persona_Badge: “Primary: AI Worker”
 - Creative Brief: Title text-6xl with cyan accent color (data role). Icon inline with title. Persona badge as small chip/pill. White text.
- Block 2:
 - Block Type: Text
 - Placement: Left column, top half
 - Component Schema: “Data Architecture Skills Card”
 - Content:
 - Section_Title: “Data Architecture for AI/ML”
 - Icon: “Layers”
 - Skill_Areas: {
 - “Feature Engineering Infrastructure”: [
 - “Feature stores (Feast, Tecton, SageMaker Feature Store)”,
 - “Feature versioning and lineage tracking”,
 - “Online vs offline feature computation”,
 - “Feature discovery and reusability”
 -],
 - “Data Quality & Validation”: [
 - “Automated data quality checks”,
 - “Schema validation and evolution”,
 - “Data profiling and anomaly detection”,
 - “Great Expectations, Deequ frameworks”
 -],
 - “ML Data Stores”: [
 - “Vector databases (Pinecone, Weaviate, pgvector)”,
 - “Time-series optimized storage”,
 - “Graph databases for relationships”,

"Data lake vs data warehouse for ML"

]

}

- Creative Brief: Card with slate-800 background, cyan accent. Three skill areas stacked. Bold area titles (text-xl), bullets in text-lg. Clear visual separation with adequate spacing.

- Block 3:

- Block Type: Text

- Placement: Left column, bottom half

- Component Schema: "Pipeline Engineering Skills Card"

- Content:

Section_Title: "ML Data Pipeline Engineering"

Icon: "GitBranch"

Skill_Areas: {

"Pipeline Orchestration": [

"Workflow orchestration (Airflow, Prefect, Kubeflow)",

"DAG design for ML workflows",

"Dependency management and scheduling",

"Pipeline monitoring and alerting"

],

"Data Processing": [

"Batch vs streaming processing for ML",

"Distributed data processing (Spark, Dask)",

"Data transformation for model training",

"ETL/ELT patterns for ML data"

],

"Data Versioning": [

"Dataset versioning (DVC, LakeFS)",

"Reproducible data snapshots",

"Data provenance and lineage",

"Training/validation/test splits management"

]

}

- Creative Brief: Matching card style to Block 2. Cyan accent throughout. Same structure and typography.

- Block 4:

- Block Type: Text

- Placement: Right column, top area (growing to fill space)

- Component Schema: "Data Pipeline Flow Diagram"

- Content:

Diagram_Title: "ML Data Pipeline Architecture"

Flow_Stages: [

{

"Stage": "Data Sources",

"Icon": "Database",

"Examples": ["Databases", "APIs", "Streaming"]

},

```
{
  "Stage": "Ingestion & Validation",
  "Icon": "Download",
  "Examples": ["Schema checks", "Quality gates"]
},
{
  "Stage": "Feature Engineering",
  "Icon": "Cpu",
  "Examples": ["Transform", "Aggregate", "Store"]
},
{
  "Stage": "Feature Store",
  "Icon": "Archive",
  "Examples": ["Online", "Offline", "Registry"]
},
{
  "Stage": "Model Training/Inference",
  "Icon": "Brain",
  "Examples": ["Consume features", "Log metrics"]
}
]
```

- Creative Brief: Vertical flow diagram with slate-700 background. Five stages connected by arrows. Each stage is a card with icon (cyan), bold stage name (text-xl), and small examples list (text-base). Clean, technical flow visualization. Arrows or connecting lines between stages.

- Block 5:

- Block Type: Text
- Placement: Right column, bottom area (sized to content)
- Component Schema: “Tools & Technologies Card”
- Content:

Section_Title: “Key Tools & Technologies”
 Icon: “Tool”
 Categories: {
 “Orchestration”: [“Apache Airflow”, “Prefect”, “Kubeflow Pipelines”, “AWS Step Functions”],
 “Processing”: [“Apache Spark”, “Dask”, “Ray”, “Flink”],
 “Feature Stores”: [“Feast”, “Tecton”, “AWS Feature Store”, “Databricks Feature Store”],
 “Versioning”: [“DVC”, “LakeFS”, “Pachyderm”]
 }

- Creative Brief: Compact card with slate-700 background, cyan accent. Four categories in a mini-grid (2x2) or stacked. Category names bold (text-lg), tools in smaller text (text-base). Very scanable.

- Block 6:

- Block Type: Text
- Placement: Footer
- Component Schema: “Simple Footer”

- Content:

Footer_Text: "Training estimate: 60-80 hours | Prerequisites: Strong data engineering fundamentals, SQL, Python"

- Creative Brief: Centered, text-base, opacity-70.